

CS536 Final Project

Josh Reese, Nadya Ortiz

December 10, 2013

0.1 Goals

In this project our aim was to mimic distributed file transfers in a simulated data center topology. From these transfers we look to measure the time it takes for all receivers to complete using both TCP and MPTCP (with varying degrees of subflows). Our interest here is under what scenarios MPTCP is useful and when does it not provide any benefit. In addition to this, when we find that MPTCP is useful how much difference does it make in these scenarios? What we expect to show here is that in a data center environment where there is a large amount of traffic and / or congestion we can see benefits to allowing nodes to send across different paths using MPTCP.

0.2 Motivation

in data centers there are often multiple paths between two nodes. If the network gets congested under TCP this can cause the amount of time a transfer will take to suffer greatly. This congestion can easily occur in TCP if there are a large number of senders and receivers and the senders have chosen routes that overlap, or are the same. However, if a sender has multiple paths to each receiver it can utilize those to help avoid congestion in the network and reduce the amount of time to transfer the file. The ideas behind MPTCP can be found in [3]. We used this work as the initial building block for our idea. We then came across some work done at Stanford which investigated various properties of a data center topology under TCP and MPTCP [4] which provided us with a working version of MPTCP for Mininet and a topology to begin working with.

0.3 Original Results

In this experiment [4] the authors created an implementation of MPTCP inside of Mininet which they used to test and verify the results from [3]. To do this they use an ECMP hashed routing protocol which is implemented in the riplpox controller inside of Mininet. To run these experiments, which include throughput, RTT, CPU utilization, and queue size, they created a Fat Tree topology in which to perform each test. A diagram of which can be found in Figure A.8. From this work we took the Fat Tree topology and used this to run our tests as well. We also use the MPTCP implementation from here. Code for this project and the mptcp setup can be found at https://github.com/bocon13/mptcp_setup and http://github.com/bocon13/datacenter_mptcp. This work helped provide us with a framework for testing as well as a topology to use in the final analysis.

0.4 Our Goals

The first goal of our project was to obtain the code from the two projects discussed in Section 0.3 and set up a working version of their MPTCP implementation using Mininet 2.0.0. The initial configuration required is an installation of their kernel update and a verification to see if it is working correctly. This takes the form of an N-switch 2-host topology, which we later used for the initial phase of testing. Followed by this the authors provide a working Fat Tree topology. It was our goal to use both of these topologies for testing. From here we would try to implement a third data center topology of our choosing and gather results from that as well.

Once these topologies were set up and our tests able to run our main goal was to investigate the impact TCP and MPTCP have on varying workloads in these environments. To accomplish this we constructed a very simple sender / receiver scenario which transfers a file (or files) from one or more senders to one or more receivers. In this way we were able to vary the amount of traffic introduced into the network.

0.5 Results

0.5.1 Setup

To conduct our experiments we set up a simple sender and receiver scenario which could be scaled to multiple senders and multiple receivers. Each sender will simply read in a file of size 50M and send a portion of it to each receiver present. Meaning if this is a 1 sender 1 receiver workload the sender will receiver all 50M of data. If we are using a 1 sender 2 receiver workload then each receiver will receiver 25M of the original data. Each sender will send 500 lines of its file at a time (roughly 64KB of data) and the receivers buffer has been set to 65536 to accommodate this. In addition to this we use the N-switch 2-host topology from https://github.com/bocon13/mptcp_setup and the Fat Tree topology from [4] to run our workloads. All experiments were conducted on Mininet 2.0.0 using the 3.5.0-89-mptcp patched kernel from https://github.com/bocon13/mptcp_setup. The machine used to conduct these experiments was a MacBook Air running OS X 9 with a 1.3 GHz core i5 processor and 4GB 1600MHz DDR3 memory. A detailed explanation of the implementation as well as some specifics about how the workloads were created can be found in Section 0.8.

0.5.2 N-switch 2-host

For the first series of tests we examined the N-switch 2-host topology with switches ranging from 2 to 5 and subflows ranging from 1 to the number of switches present. In these tests we have 1 sender and 1 receiver connected to variable amounts of switches and we examine the time it takes to transfer a file of size 50M between the two.

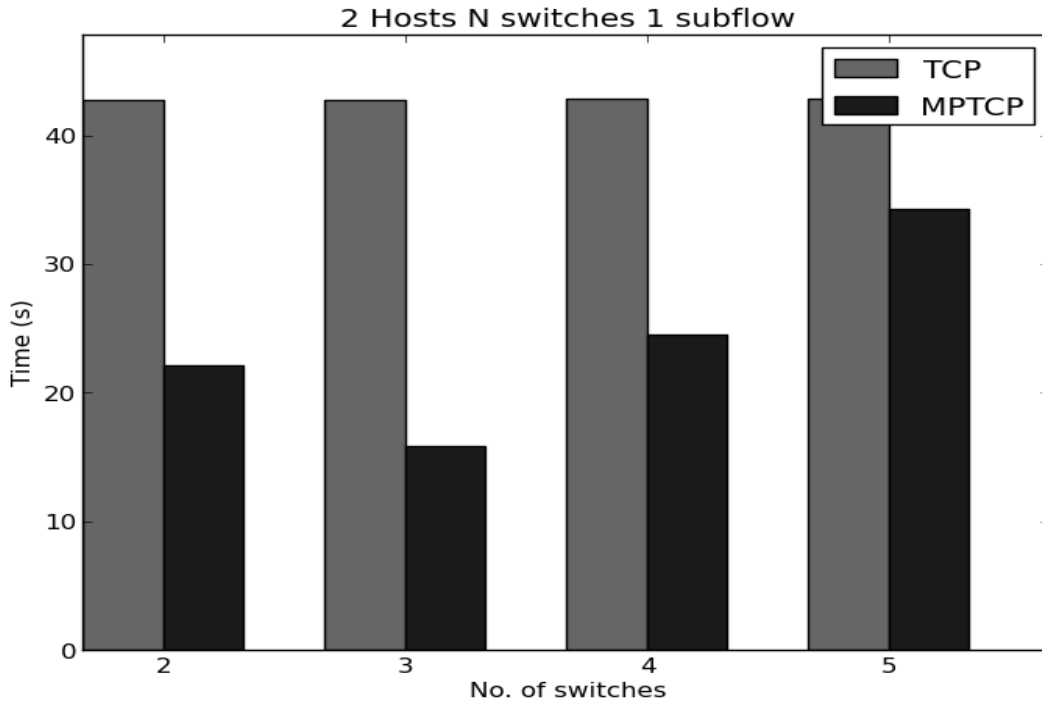


Figure 1: Time vs number of switches for 1 subflow.

0.5.3 Fat Tree Topology (k=4)

In the second series of testing we ran our sender / receiver implementation on the Fat Tree topology (of k=4) with various number of senders and receivers. One point to note is that we found cases where the limitations in our system, either in hardware or with Mininet, prevented us from running some tests. This occurs when there are too many flows coupled with multiple senders and receivers. Specifically we noticed this cases of more than 5 senders, 6 receivers, and 7 to 8 subflows. Similar issues were discovered in [4] regarding CPU limitations on their system. In these cases we don't report results for these tests because the results are unpredictable and thus don't reliably represent the experiment. We represent these results in individual charts by number of receivers (1 to 6) with time to transfer a file (in seconds) on the y-axis and number of subflows on the x-axis. Each line in the graph corresponds to a different number of senders. In all of these graphs the point on the y-axis (1 subflow) represents the time to transfer the file(s) under TCP while the remaining points represent the transfer time with MPTCP and the corresponding number of subflows. These charts can be found in Appendix A.

0.6 Interpretation of Results

In the tests presented in Section 0.5.2, specifically Figure A.1, we can examine the graphs that test number of flows vs. time. From these graphs we can see that adding 1 additional subflow can drastically decrease the transfer time between the 2 hosts. However, when the number of flows increases to above 1 the implementation appears to revert back to TCP. This is something we couldn't figure out and spent some time examining but never came to any conclusions. If we examine Figure 1 we can see a much clearer picture of the benefits of using MPTCP in this topology. In this graph we show that with more paths between the sender and receiver MPTCP can effectively cut the transfer time in half. Another thing to note is that as the switches increase past 3 the time to transfer using MPTCP tends to increase. Our intuition regarding the cause of this suggests two possible scenarios: the first is that on the senders side the number of paths available is increasing and thus the sender will have to monitor more paths according to the MPTCP protocol. On the other side of the transfer, the receiver is likely receiving more information simultaneously from all the switches and the threads performing the receiving are forced to switch back and forth rapidly.

If we now look at Figures A.2-A.7 we can see some interesting results. In Figure A.2 we see there is really no advantage to using MPTCP regardless of the number of senders. Our assumption here is that the traffic is congested in the switch located directly above the host in this topology since all senders are sending to only one location. In graphs past Figure A.3 we no longer use subflows above 6 due to the limitation previously mentioned. Starting with this plot we can begin to see some advantages to using MPTCP (specifically here for eight senders). As the number of receivers increases we can more clearly see the advantages to using MPTCP over TCP. In some cases decreasing the transfer time by approximately 30%. These results appear to support our initial assumptions very nicely. We can see that as the number of receivers increase, thus giving the senders more options regarding possible paths, we are able to begin to see the benefits of using MPTCP. Table A.1 provides a concise presentation of the times taken for 3, 4, 5 and 6 receivers and the maximum amount of senders our system could handle for those receivers. Here we report specific times for TCP and MPTCP with 2, 3, 4 and 5 subflows. From this table we can clearly see that in the 4 and 5 receiver tests creating even 1 addition subflow can cause the transfer time to drastically decrease (79s to 53s for 5 receivers).

0.7 Challenges

The first obstacle in this project was getting a working version of Mininet with MPTCP up and running. From here we had to understand the code used in [4] which included two fairly extensive python files controlling the routing in the network and the design of the topologies. Each of these files contained dependencies on various Mininet components requiring us to dig deep into the internals of the system to try to understand what was going on. When we had our initial system running and were performing our early tests we noticed we couldn't obtain any performance increase from using MPTCP. This caused a bit of a panic and we came to the conclusion that the edge switches inside the Fat Tree were causing too much of a bottleneck for our file transfers. From here we developed an entirely different topology, a Dual Homed Tree, which we found in [3][2]. We struggled with the Dual Homed Tree, shown in Figure A.9

for some time, changing naming conventions and interfaces and struggling with the riplpox controller inside Mininet. We finally discovered from [1] that our routing might be wrong. This gave us key insights into implementing a functioning version of the Dual Homed Tree. In the end however, we were unable to obtain any results from this topology due to limitations in connecting multiple interfaces from a host to separate switches inside of Mininet. Even after we had this new topology set up we spent a lot of time examining the output of the controller to see if hosts were actually taking multiple paths in their file transfers. One additional issue we encountered during the development of the Dual Homed Tree was a situation where Mininet entered a state in which it wasn't able to cleanup all the links and switches it had created. We investigated through searching the documentation and FAQs and posted onto the mailing list. In the end we found out it was a kernel bug.

During the testing process we encountered a very large number of errors. The most daunting of these was the amount of time required to perform these tests. Due to the length of time for some of the transfers we were unable to perform multiple executions of the test to generate more accurate results. Once we started testing and really stressing the system we noticed that on occasion Mininet would stop responding and a transfer could continue indefinitely. This occurred as the receivers increased and in flows of size seven and eight. At this point we would need to reboot the system and begin the test at the last stable checkpoint. This explains some of the unexpected points of data in our plots. Some of these challenges are also mentioned in [4] and are intrinsic to Mininet itself. In addition to this we found that sometimes even for tests with low stress on the system we would generate unexpected results. This could be due to any number of factors, including issues with Mininet not being properly cleaned up from previous tests.

0.8 Implementation

Code and instructions for recreating the tests from this project can be found at: <https://github.com/jreeseue/mptcp>. A dataset is provided under the data/covtype directory to use for testing. There are several files of note in this implementation.

- **sender.py:** This file contains the code for a node to send its data. A sender takes an ID (used to read its file), a chunk size (unused), a dataset (unused but could be used to support different files in the future), and a list of IP addresses to send to. A sender will transfer its file, or a portion of its file, to each receiver in its list of IPs. Each sender will send 500 lines of data (approximately 64KB in this dataset) at a time until it has finished its transfer.
- **receiver.py:** This file contains the code for a node to receive data from one or more senders. A receiver takes an ID (used to save its data), a chunk size (unused), the number of receivers, the number of senders, and the name of the dataset as parameters. Each receiver will listen for connections and create a thread for each incoming connection. These threads will correspond to each sender in the test and will receive data until all the senders have completed their transfers. Each receiver has a buffer of size 65536 to correspond to the approximate 64KB each sender will be sending in each transmission. Once all senders are finished all receivers will write their received files to the disk.
- **test_ft.py:** This file contains the code to test the implementation of the Fat Tree topology using the sender and receiver code previously described. Here we set up Mininet and attach it to the riplpox controller for routing, construct custom links, and create the mappings between senders and receivers. In an attempt to distribute the workload throughout the network we ensure that senders and receivers are never located in the same pod. To do this we assign pods 0 and 2 to senders and pods 1 and 3 to receivers. When a test is run senders and receivers are chosen at random from these pods.
- **dctopo.py:** This is a modified Fat Tree topology file from https://github.com/bocon13/datacenter_mptcp. The only changes to this are the Dual Homed Tree topology we've added at the end (which currently doesn't function). This is the topology used for testing the Fat Tree.
- **run_test.py:** A script to run the workload of tests for the Fat Tree. Senders 1 to 8, receivers 1 to 8, and flows 2 to 6. Not all of these tests were able to run on our system.

- **run_switch.py:** A script to run the workload of tests for the N-switch 2-host setup. Switches ranging from 2 to 5.

0.9 Conclusion

From examining the work presented here we can see very clear benefits to using MPTCP in a data center which has a high volume of traffic and congestion. As we demonstrated with the increasing amount of receivers we see also an increasing benefit to using MPTCP to transfer a file. In addition to this we can see that for all the tests we performed when the number of subflows reached four the time to transfer stopped decreasing. This suggests either our system wasn't able to properly handle flows above four or in this test set up four subflows may be optimal in most cases. To conclude, the benefits to using MPTCP in a higher volume data center may be very significant.

0.10 Future Work

There are several possible ways in which this work could be expanded. This first, and most obvious of which, would be to implement other data center topologies inside of Mininet and perform a comparison between those and the Fat Tree used here. This could provide us with an interesting comparison between how the physics layout of a data center might have an impact on the effectiveness of using MPTCP. While we tried to implement the Dual Homed Tree topology, there wasn't enough time to get this running and perform the tests. Two other topologies we examined were VL2 and the B-Cube. Either of these two might prove to be an interesting comparison to the Fat Tree, or each other. In addition there are various parameters inside of Mininet that could be tweaked for measurements. These include the bandwidth of links and size of packets allowed in the queue of each switch (measured in [4]). Once again these parameters might provide interesting results when moving from TCP to MPTCP. As a modification to the code we produced here an interesting change might be in the amount of data a sender will send at a time and the size of a receivers buffer. The values we set were initially fairly large (64KB) and it might prove interesting to measure the impact different buffer sizes (and send amounts) might have in these environments.

Appendix A

Plots

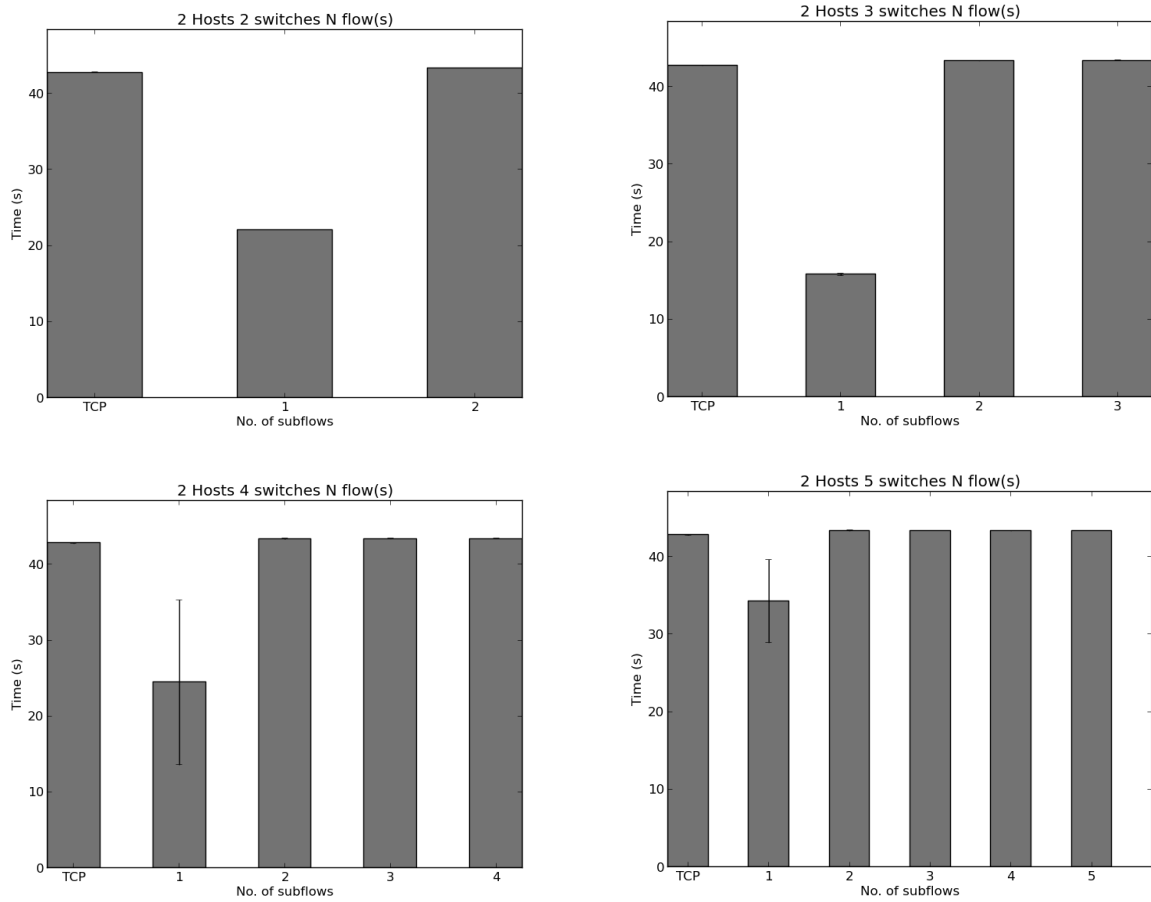


Figure A.1: Time vs. Number of flows for switches ranging from 2 to 5.

# Recv	# Sndr	TCP	1 Flow	2 Flows	3 Flows	4 Flows	5 Flows
3	5	87	77	80	79	77	81
4	5	77	72	64	65	65	65
5	5	79	53	50	50	52	56
6	5	62	53	53	50	53	50

Table A.1: Number of receivers with the maximum number of senders that could be used during testing. Times reported are for TCP and the various number of subflows.

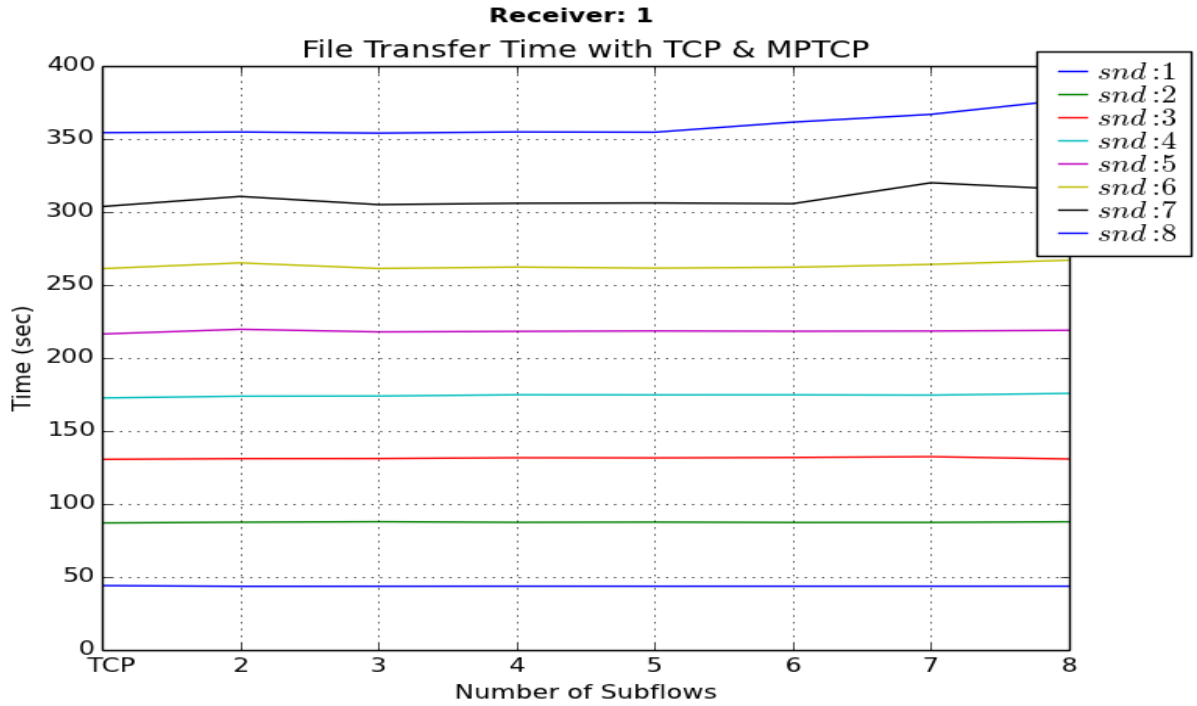


Figure A.2: 1 receiver multiple senders.

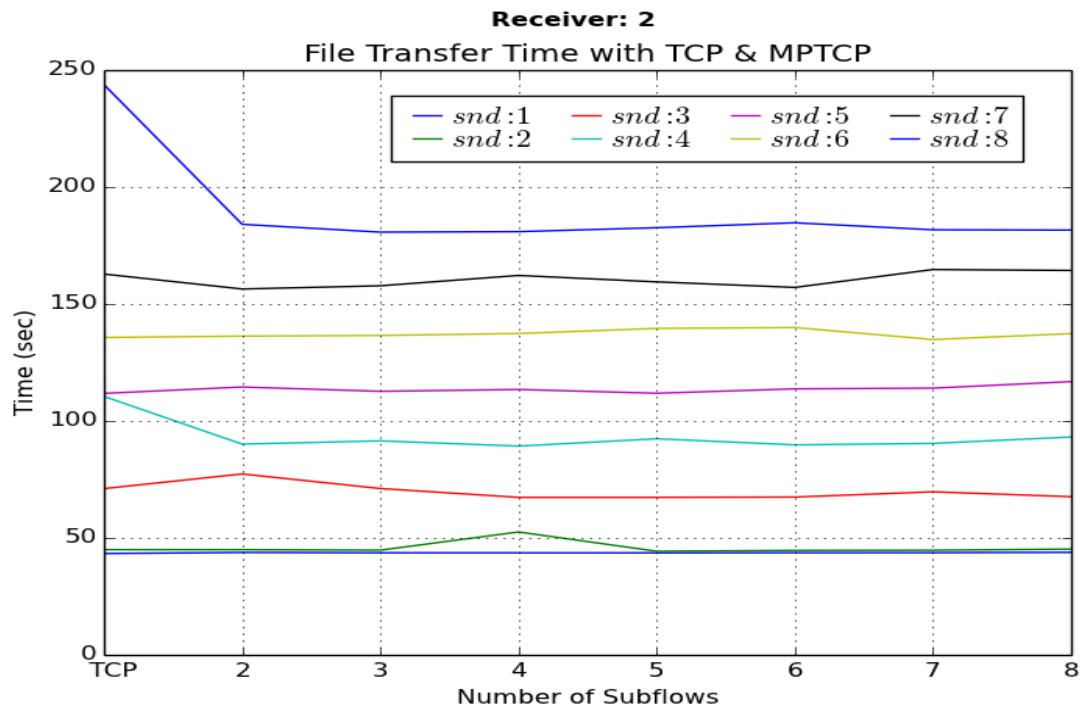


Figure A.3: 2 receivers multiple senders.




Figure A.4: 3 receivers multiple senders.




Figure A.5: 4 receivers multiple senders.

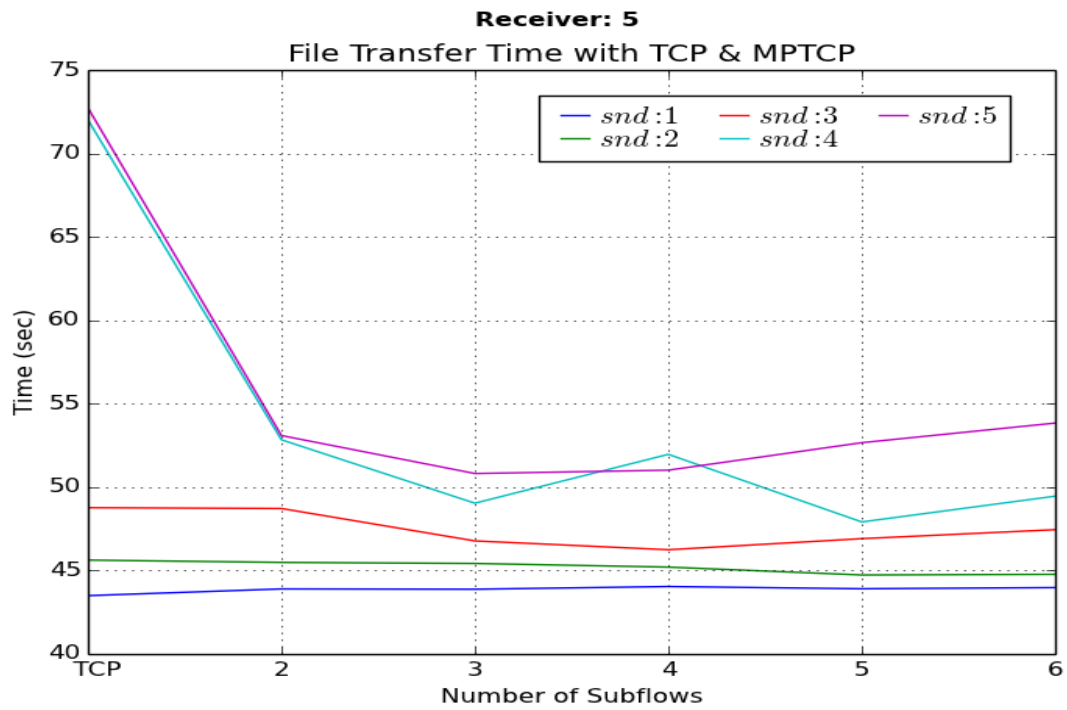


Figure A.6: 5 receivers multiple senders.




Figure A.7: 6 receivers multiple senders.




Figure A.8: Fat Tree Topology with $k=4$.




Figure A.9: Dual Homed Tree Topology with $k=4$.

Bibliography

- [1] Mohammad Al-Fares, Alexander Loukissas, and Amin Vahdat. A scalable, commodity data center network architecture. *SIGCOMM Comput. Commun. Rev.*, 38(4):63–74, August 2008.
- [2] Costin Raiciu. http://changeofelia.info.ucl.ac.be/pmwiki/uploads/SummerSchool/Program/session_010.pdf.
- [3] Costin Raiciu, Sebastien Barre, Christopher Pluntke, Adam Greenhalgh, Damon Wischik, and Mark Handley. Improving datacenter performance and robustness with multipath tcp. In *Proceedings of the ACM SIGCOMM 2011 Conference, SIGCOMM '11*, pages 266–277, New York, NY, USA, 2011. ACM.
- [4] Brian O'Connor Ryan Wilson. <http://reproducingnetworkresearch.wordpress.com/2013/03/13/cs244-13-improving-datacenter-performance-and-robustness-with-multipath-tcp/>, 2013.